

S&DS 265 / 565
Introductory Machine Learning

Course Wrap Up

December 7

Yale

Endgame

- Quiz 6 open today at 1pm; closes Friday at 6pm
- Final exam: Friday, Dec 15 at 2pm in SSS 114
- Practice exams posted
- Review sessions:
 - ▶ Regina: Wednesday (Dec 13): 5-7pm Bass L73
 - ▶ Kaylee: Thursday: (Dec 14) 5-7pm KT 211
 - ▶ Hannah: Thursday: (Dec 14) 7-9pm KT 211

Last unit: Language/Sequence models

- Generative process, any sequence (of words, characters, stock prices, nucleotides...) is assigned a probability

$$p(x_1, \dots, x_n)$$

which can be factored as

$$p(x_1, \dots, x_n) = p(x_1)p(x_2 | x_1) \dots p(x_n | x_1, \dots, x_{n-1})$$

Transformers

The current state-of-the-art is based on *transformers*

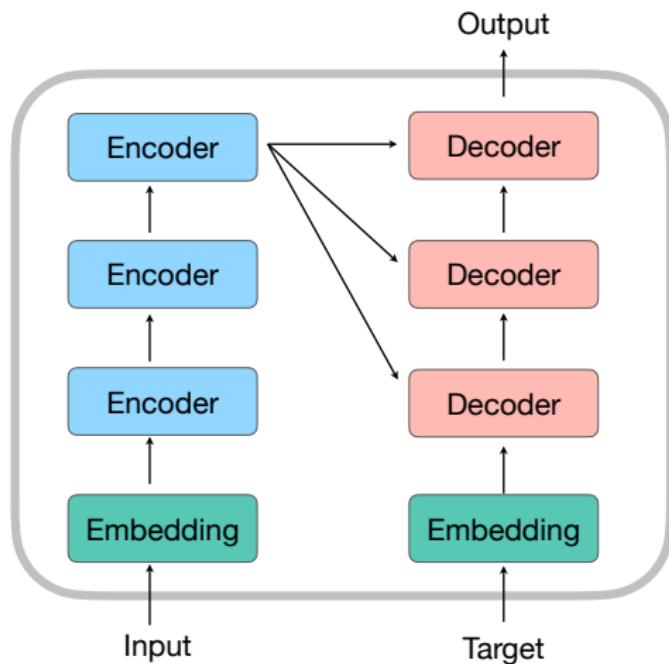
- Attention is the key ingredient
- Rather than processing sequences word-by-word, transformers handle larger chunks of text at once
- Incorporate “interactions” between words and hidden states

Transformer architecture

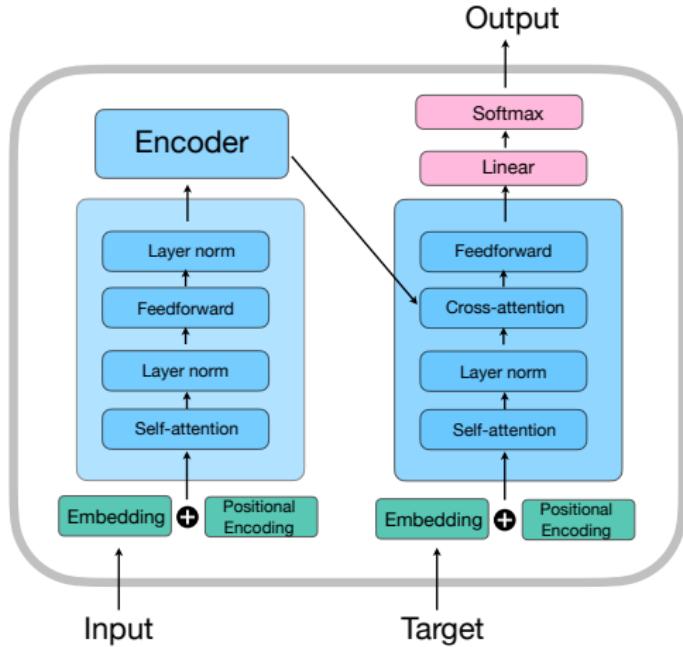
A Transformer is a seq2seq model based on encoder and decoder modules.

Transformers are powerful alternatives to RNNs that transform the encoder/decoder states using (multi-head) attention mechanisms.

Transformer architecture

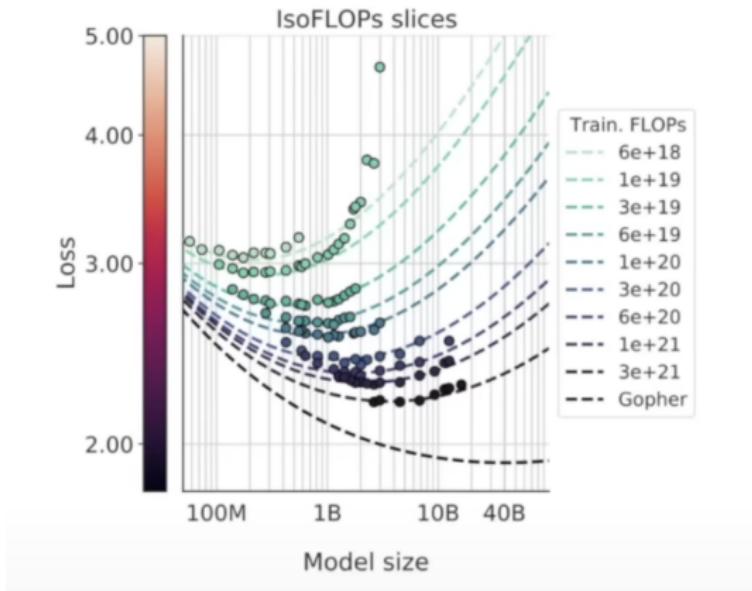


Transformer architecture



Two encoder layers and one decoder layer

LLM scaling laws: Bigger is better

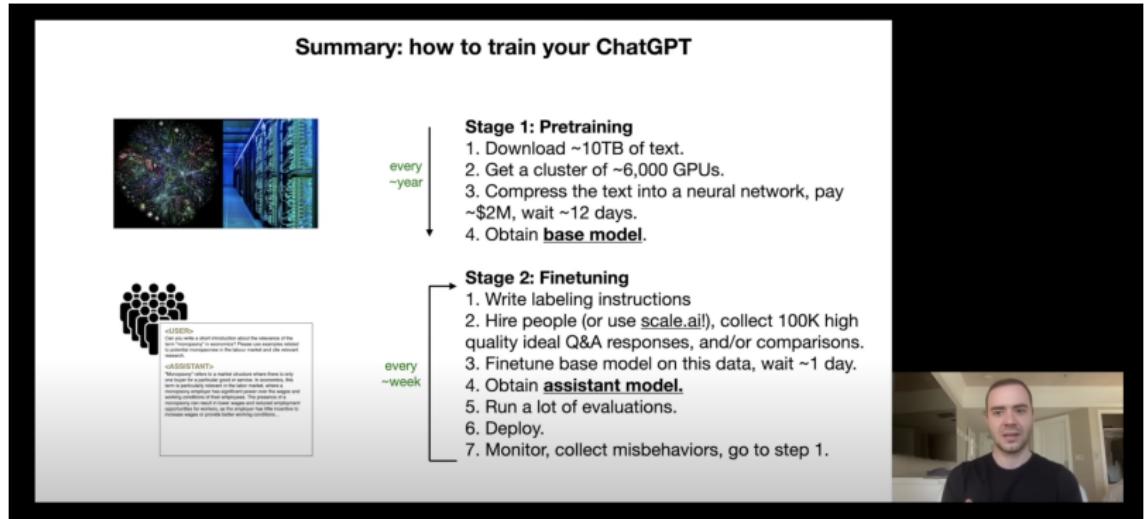


https://www.youtube.com/watch?v=zjkBMFhNj_g&t=25m40s

Finetuning (Training an LLM Decoder)

- An incredible amount of knowledge is stored implicitly in the weights of the transformer
- To be made useful, supervised learning and reinforcement learning are used to teach the model how to respond

Finetuning



https://www.youtube.com/watch?v=zjkBMFhNj_g&t=14m19s

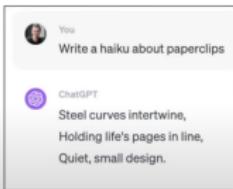
Finetuning

RLHF

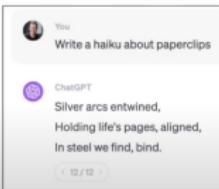
The second kind of label: comparisons

It is often much easier to compare Answers instead of writing Answers.

Simple example: it's much easier to spot a good haiku than it is to generate one:

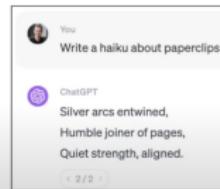


ChatGPT
Steel curves intertwine,
Holding life's pages in line,
Quiet, small design.



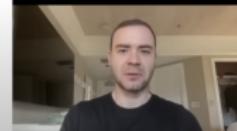
ChatGPT
Silver arcs entwined,
Holding life's pages, aligned.
In steel we find, bind.

12 / 12



ChatGPT
Silver arcs entwined,
Humble joiner of pages,
Quiet strength, aligned.

13 / 12



https://www.youtube.com/embed/zjkBMFhNj_g?start=1277&end=1340

Step 1

Collect demonstration data and train a supervised policy.

A prompt is sampled from our prompt dataset.

Explain reinforcement learning to a 6 year old.



We give treats and punishments to teach...



SFT



¶¶¶

A labeler demonstrates the desired output behavior.

This data is used to fine-tune GPT-3.5 with supervised learning.

Step 2

Collect comparison data and train a reward model.

A prompt and several model outputs are sampled.

Explain reinforcement learning to a 6 year old.

A
It's reinforcement learning, the agent...

B
Dolphin research...

C
In machine learning...

D
We give treats and punishments to teach...



D > C > A > B

A labeler ranks the outputs from best to worst.

This data is used to train our reward model.

RM



¶¶¶

This data is used to train our reward model.

Step 3

Optimize a policy against the reward model using the PPO reinforcement learning algorithm.

A new prompt is sampled from the dataset.

Write a story about otters.



Once upon a time...

PPO

RM

r_k

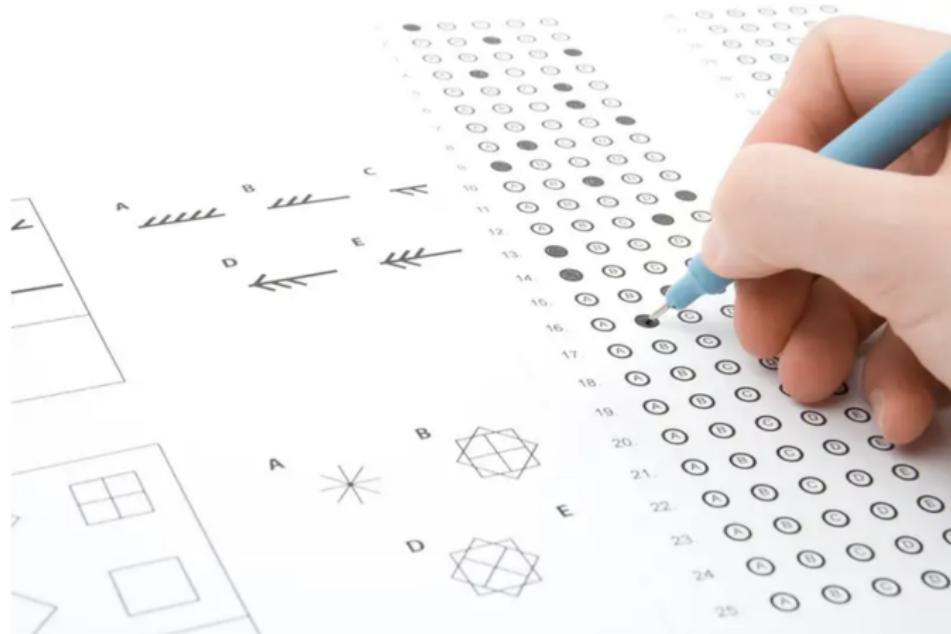
The PPO model is initialized from the supervised policy.

The policy generates an output.

The reward model calculates a reward for the output.

The reward is used to update the policy using PPO.

What's next?: Fast learning, slow thinking



AI struggles with IQ tests
Panther Media GmbH / Alamy Stock Photo

Always news—on both sides

The
Economist

≡ Menu | Weekly edition | The world in brief | Q, Search ▾

Science and technology | Crystal balls

A Google AI has discovered 2.2m materials unknown to science

Zillions of possible crystals exist. AI can help catalogue them



Just one of trillions IMAGE: SCIENCE PHOTO LIBRARY

Nov 29th 2023

Save

Share

Give

Always news—on both sides

The
Economist

≡ Menu

Weekly edition

The world in brief

Q Search ▾

Business | Of evils and evals

The world wants to regulate AI, but does not quite know how

There is disagreement over what is to be policed, how and by whom

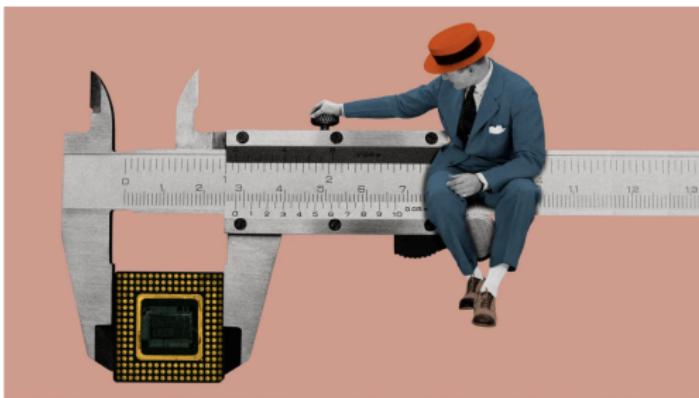


IMAGE: MARI FOUZ

Oct 24th 2023 | BLETCHLEY PARK

Save

Share

Give

Advancing ML: Something for everyone

- Designing new methods
- Applying to new domains
- Solving deep mathematical puzzles
- Tackling unique engineering challenges
- Designing interfaces
- Collecting data
- Equitable use through policy and law
- Outreach and communication to broad communities

We've covered a lot of ground!

Calendar Fall 2023

Lectures: Tuesday/Thursday 11:35-12:50pm

Davies Auditorium

Complementary readings marked ISL refer to sections in the book [An Introduction to Statistical Learning](#) (Python version, July 2023). Assignments and quizzes are posted and due on Thursday in a given week.

Week	Dates	Topics	Demos & Tutorials	Lecture Slides	Readings and Notes	Assignments & Exams
1	Aug 31	Course overview		Thu: Course overview		
2	Sept 5, 7	Python and background concepts	CO Python elements CO Covid trends	Tue: Python elements Thu: Pandas and linear regression	Data8 Chapters 3, 4, 5	Quiz 1 CO Assn 1 out
3	Sept 12, 14	Linear regression and classification	CO Covid trends (revisited) CO Classification examples	Tue: Regression concepts Thu: Classification	ISL Sections 3.1, 3.2, 3.5 Notes on regression ISL Sections 4.3, 4.4 Notes on classification	
4	Sept 19, 21	Stochastic gradient descent	CO SGD examples	Tue: Classification (continued) Thu: Stochastic gradient descent	ISL Section 6.2.2 ISL Section 10.7.2	Assn 1 in CO Assn 2 out

We've covered a lot of ground!

5	Sept 26, 28	Bias and variance, cross-validation	Bias-variance tradeoff Covid trends (revisited) California housing	Tue: Bias and variance Thu: Cross-validation	ISL Section 2.2 ISL Section 5.1	Quiz 2
6	Oct 3, 5	Tree-based methods and principal components	Trees and forests Visualizing trees PCA examples	Tue: Trees and Forests Thu: PCA	ISL Sections 8.1, 8.2 ISL Section 12.2	Assn 2 in Assn 3 out
7	Oct 10, 12	PCA and dimension reduction	PCA revisited Used for dimension reduction Word embeddings	Tue: PCA and word embeddings Thu: Embeddings and review	ISL Section 12.2	Quiz 3
8	Oct 17	Midterm exam (in class)			On Canvas: Practice midterms / Sample solns Midterm / Sample soln	

We've covered a lot of ground!

10	Oct 31, Nov 2	Topic models, introduction to neural networks	 Sanity check  Minimal neural network  Regression examples	Tue: Topic models Thu: Neural networks	ISL Sections 10.1, 10.2	Quiz 4
11	Nov 7, 9	Neural networks, reinforcement learning	 Q-learning	Tue: Neural networks Thu: Reinforcement learning	Notes on backpropagation	Assn 4 in  Assn 5 out
12	Nov 14, 16	Deep neural networks	Tensorflow playground  Autoencoder examples	Tue: Deep reinforcement learning Thu: Deep networks and autoencoders	ISL Section 10.7	Quiz 5
13	Nov 21, 23	No class, Thanksgiving break				

We've covered a lot of ground!

14	Nov 28, 30	Transformers and ChatGPT	 GPT-4 Python API	Tue: Autoencoders and transformers Thu: Transformers (continued)		
15	Dec 5, 7	Societal issues for machine learning		Tue: Panel discussion Thu: Course wrap up		Assn 5 in Quiz 6
16	Fri, Dec 15, 2pm, SSS 114	Final exam			Registrar: Final exam schedule Practice finals	

Final exam

- Final exam Friday, Dec 15, 2023 at 2pm in SSS 114
- <https://registrar.yale.edu/general-information/final-exams>
- Review sessions (see times/dates above)
- Length: About 1.5X Midterm
- Emphasis on material after midterm
- Cumulative, closed book, cheat-sheet

Your input

- Please complete a course review!
- I greatly value your comments and feedback
- Feel free to send me comments privately
- Let me know how you use and continue to learn ML!

Thank you!