



Plan for today

- Reminders
- Recap of word embeddings
- Demo notebook
- Review–AMA

Reminders

- Assn 3 is out; due October 26
- Quiz 3 posted today; open at 1pm, closes Friday at 6pm—discriminative vs. generative models, trees, bias/variance, SGD – good review for midterm!
- Midterm Tuesday, October 17, in class
- “Closed book, notes, computer...”
- $8\frac{1}{2} \times 11$ sheet of notes, handwritten double-sided
- Practice midterms posted on Canvas (with solutions)

Review sessions

Regina: Thursday, October 12, 8-9pm

Hannah: Friday, October 13, 5-6pm

Kaylee: Saturday, October 14, 5-6pm

Awni: Sunday, October 15, 2-3pm

Held in Kline Tower (KT) 211 and Zoom (see announcements)

Language models

- A language model is a way of assigning a probability to any sequence of words (or string of text)

$$p(w_1, \dots, w_n)$$

Language models

- A language model is a way of assigning a probability to any sequence of words (or string of text)

$$p(w_1, \dots, w_n)$$

- By the basic rules of conditional probability we can factor this as

$$p(w_1, \dots, w_n) = p(w_1)p(w_2 | w_1) \dots p(w_n | w_1, \dots, w_{n-1})$$

Modern language models

Suppose a computer program assigns a “score” to possible next words v :

$$s(v; \overbrace{w_1, \dots, w_n}^{\text{previous words}})$$

↑
possible next word

Modern language models

Suppose a computer program assigns a “score” to possible next words v :

$$\begin{array}{c} \text{previous words} \\ s(v; \overbrace{w_1, \dots, w_n}) \\ \uparrow \\ \text{possible next word} \end{array}$$

Can convert this to a language model by the “softmax” operation:

$$p(w \mid w_1, \dots, w_n) = \frac{\exp(s(w; w_1, \dots, w_n))}{\sum_{v \in V} \exp(s(v; w_1, \dots, w_n))}$$

Modern language models

Suppose a computer program assigns a “score” to possible next words v :

$$\begin{array}{c} \text{previous words} \\ s(v; \overbrace{w_1, \dots, w_n}) \\ \uparrow \\ \text{possible next word} \end{array}$$

Can convert this to a language model by the “softmax” operation:

$$p(w \mid w_1, \dots, w_n) = \frac{\exp(s(w; w_1, \dots, w_n))}{\sum_{v \in V} \exp(s(v; w_1, \dots, w_n))}$$

In ChatGPT, the function $s(v; w_{1:n})$ is learned on large amounts of text (unsupervised) using a type of deep neural network called a *transformer*.

Modern language models

A language model assigns a “score” to possible next words v :

$$s(v; \underbrace{w_1, \dots, w_n}_{\text{word history}})$$

Modern language models

A language model assigns a “score” to possible next words v :

$$s(v; \underbrace{w_1, \dots, w_n}_{\text{word history}})$$

Today, we'll be working with a simple case where

$$\begin{aligned} s(v; w_1, \dots, w_n) &= \beta_v^T \phi(w_1, \dots, w_n) \\ &= \beta_v^T \phi(w_n) \\ &= \phi(v)^T \phi(w_n) \end{aligned}$$

Modern language models

A language model assigns a “score” to possible next words v :

$$s(v; \underbrace{w_1, \dots, w_n}_{\text{word history}})$$

Today, we'll be working with a simple case where

$$\begin{aligned} s(v; w_1, \dots, w_n) &= \beta_v^T \phi(w_1, \dots, w_n) \\ &= \beta_v^T \phi(w_n) \\ &= \phi(v)^T \phi(w_n) \end{aligned}$$

Key intuition

- Similar words will appear with similar words
- Self-referential notion of similarity

Constructing embeddings

Language model is

$$p(w_2 | w_1) = \frac{\exp(\phi(w_2)^T \phi(w_1))}{\sum_w \exp(\phi(w)^T \phi(w_1))}.$$

Carry out stochastic gradient descent over the embedding vectors $\phi \in \mathbb{R}^d$ (where $d \approx 50$ – 500 is chosen by hand)

This is what Mikolov et al. (2014, 2015) did at Google. With a couple of twists:

Constructing embeddings

Heuristics used:

- Skip-gram: predict surrounding words from current word, rather than the next word.

Constructing embeddings

Heuristics used:

- Skip-gram: predict surrounding words from current word, rather than the next word.
- This leads to a model of nearby words $p_{\text{near}}(w_2 | w_1)$.

Constructing embeddings

Heuristics used:

- Skip-gram: predict surrounding words from current word, rather than the next word.
- This leads to a model of nearby words $p_{\text{near}}(w_2 | w_1)$.
- Second is computational, to avoid bottleneck in computing the denominator in the softmax.

GloVe

Shortly after, a group at Stanford group introduced a variant called “GloVe”

- Based on a type of regression model
- More scalable with SGD

Using PCA

A closely related approach is to use PCA of pointwise mutual information (PMI):

- Form $V \times V$ matrix of pointwise mutual information values

$$\log \left(\frac{p_{\text{near}}(w_1, w_2)}{p(w_1)p(w_2)} \right)$$

- Compute top k eigenvectors ϕ_1, \dots, ϕ_k
- For each word w , define embedding as

$$\phi(w) \equiv (\phi_{1w}, \phi_{2w}, \dots, \phi_{kw})^T$$

Analogies

Leads to vector representations of words with interesting properties.

For example, analogies:

`king is to man as ? is to woman`

Analogies

Leads to vector representations of words with interesting properties.

For example, analogies:

king is to man as ? is to woman

Paris is to France as ? is to Germany

Analogies

Leads to vector representations of words with interesting properties.

For example, analogies:

king is to man as ? is to woman

Paris is to France as ? is to Germany

$$\phi(\text{king}) - \phi(\text{man}) \stackrel{?}{\approx} \phi(\text{queen}) - \phi(\text{woman})$$

$$\hat{w} = \arg \min_w \|\phi(\text{king}) - \phi(\text{man}) + \phi(\text{woman}) - \phi(w)\|^2$$

Does $\hat{w} = \text{queen}$?

Learned Analogies

Table 8: *Examples of the word pair relationships, using the best word vectors from Table 4 (Skip-gram model trained on 783M words with 300 dimensionality).*

Relationship	Example 1	Example 2	Example 3
France - Paris	Italy: Rome	Japan: Tokyo	Florida: Tallahassee
big - bigger	small: larger	cold: colder	quick: quicker
Miami - Florida	Baltimore: Maryland	Dallas: Texas	Kona: Hawaii
Einstein - scientist	Messi: midfielder	Mozart: violinist	Picasso: painter
Sarkozy - France	Berlusconi: Italy	Merkel: Germany	Koizumi: Japan
copper - Cu	zinc: Zn	gold: Au	uranium: plutonium
Berlusconi - Silvio	Sarkozy: Nicolas	Putin: Medvedev	Obama: Barack
Microsoft - Windows	Google: Android	IBM: Linux	Apple: iPhone
Microsoft - Ballmer	Google: Yahoo	IBM: McNealy	Apple: Jobs
Japan - sushi	Germany: bratwurst	France: tapas	USA: pizza

Evaluation Analogies

Type of relationship	Word Pair 1		Word Pair 2	
Common capital city	Athens	Greece	Oslo	Norway
All capital cities	Astana	Kazakhstan	Harare	Zimbabwe
Currency	Angola	kwanza	Iran	rial
City-in-state	Chicago	Illinois	Stockton	California
Man-Woman	brother	sister	grandson	granddaughter
Adjective to adverb	apparent	apparently	rapid	rapidly
Opposite	possibly	impossibly	ethical	unethical
Comparative	great	greater	tough	tougher
Superlative	easy	easiest	lucky	luckiest
Present Participle	think	thinking	read	reading
Nationality adjective	Switzerland	Swiss	Cambodia	Cambodian
Past tense	walking	walked	swimming	swam
Plural nouns	mouse	mice	dollar	dollars
Plural verbs	work	works	speak	speaks

Notebook

Let's go to the Python notebook!

Embedding / Visualization Examples

[WebVectors](#)[Similar words](#)[Visualizations](#)[Calculator](#)[2D text](#)[Miscellaneous](#)[Models](#)[About](#)

WebVectors: word embeddings online

"You shall know a word by the company it keeps." (Firth 1957)

Enter a word to produce a list of its 10 nearest semantic associates.
English Wikipedia model will be used; for other models, visit [Similar Words](#) tab.

Semantic associates for **platypus** (computed on [English Wikipedia](#))

Word frequency

☒ High ☒ Medium ☐ Low

1. marsupial 0.642
2. crocodile 0.605
3. kangaroo 0.595
4. turtle 0.595
5. iguana 0.589
6. frog 0.573
7. penguin 0.572
8. wallaby 0.570
9. alligator 0.569
10. opossum 0.568



Similarity threshold ☐ Show tags

• We show only the associates of the same part of speech as your query. All associates can be found at the [Similar Words](#) tab.

<http://vectors.nlp1.eu/explore/embeddings/en/>

Many uses

species2vec: A novel method for species representation

 Boyan Angelov

doi: <https://doi.org/10.1101/461996>

This article is a preprint and has not been certified by peer review [what does this mean?].

Abstract

Full Text

Info/History

Metrics

 Preview PDF

Abstract

Word embeddings are omnipresent in Natural Language Processing (NLP) tasks. The same technology which defines words by their context can also define biological species. This study showcases this new method - species embedding (species2vec). By proximity sorting of 6761594 mammal observations from the whole world (2862 different species), we are able to create a training corpus for the skip-gram model. The resulting species embeddings are tested in an environmental classification task. The classifier performance confirms the utility of those embeddings in preserving the relationships between species, and also being representative of species consortia in an environment.

Visualisation

```
In [10]: m = gensim.models.KeyedVectors.load_word2vec_format('reptilia.vec')
```

```
In [11]: len(m.vocab)
```

```
Out[11]: 7397
```

```
In [15]: m.most_similar(u'Alligator_mississippiensis')
```

```
Out[15]: [(u'Sternotherus_bonevalleyensis', 0.8425856828689575),  
(u'Apalone_ferox', 0.8147842884063721),  
(u'Macrochelys_suwanniensis', 0.8063992261886597),  
(u'Deirochelys_reticularia', 0.7871163487434387),  
(u'Terrapene_putnami', 0.7841686010360718),  
(u'Chelydra_floridana', 0.7829421758651733),  
(u'Alligator_mefferdi', 0.7742743492126465),  
(u'Macrochelys_temminckii', 0.7682404518127441),  
(u'Trachemys_inflata', 0.7563525438308716),  
(u'Deirochelys_carri', 0.755811333656311)]
```

```
In [16]: %matplotlib inline
```

```
def tsne_plot(model):  
    "Creates and TSNE model and plots it"  
    labels = []
```

Summary: Word embeddings

- Word embeddings are vector representations of words, learned from cooccurrence statistics
- The models can be built using language modeling (or regression or PCA)
- Surprising semantic relations are encoded in linear relations—for example, analogies
- Embeddings are the “ground floor” representations in ChatGPT

Week	Dates	Topics	Demos & Tutorials	Lecture Slides	Readings and Notes	Assignments & Exams
1	Aug 31	Course overview		Thu: Course overview		
2	Sept 5, 7	Python and background concepts	Python elements Covid trends	Tue: Python elements Thu: Pandas and linear regression	Data8 Chapters 3, 4, 5	Quiz 1 Assn 1 out
3	Sept 12, 14	Linear regression and classification	Covid trends (revisited) Classification examples	Tue: Regression concepts Thu: Classification	ISL Sections 3.1, 3.2, 3.5 Notes on regression ISL Sections 4.3, 4.4 Notes on classification	
4	Sept 19, 21	Stochastic gradient descent	SGD examples	Tue: Classification (continued) Thu: Stochastic gradient descent	ISL Section 6.2.2 ISL Section 10.7.2	Assn 1 in Assn 2 out
5	Sept 26, 28	Bias and variance, cross-validation	Bias-variance tradeoff Covid trends (revisited) California housing	Tue: Bias and variance Thu: Cross-validation	ISL Section 2.2 ISL Section 5.1	Quiz 2
6	Oct 3, 5	Tree-based methods and principal components	Trees and forests Visualizing trees PCA examples	Tue: Trees and Forests Thu: PCA	ISL Sections 8.1, 8.2 ISL Section 12.2	Assn 2 in Assn 3 out
7	Oct 10, 12	PCA and dimension reduction	PCA revisited Used for dimension reduction Word embeddings	Tue: PCA and word embeddings Thu: Review	ISL Section 12.2	Quiz 3

“Ask Me Anything” (AMA)